



AB Testing for Process Versions with Contextual Multi-armed Bandit Algorithms

Suhrid Satyal (presenter),

Ingo Weber, Hye-young Paik, Claudio Di Ciccio, Jan Mendling

June 2018

www.csiro.au

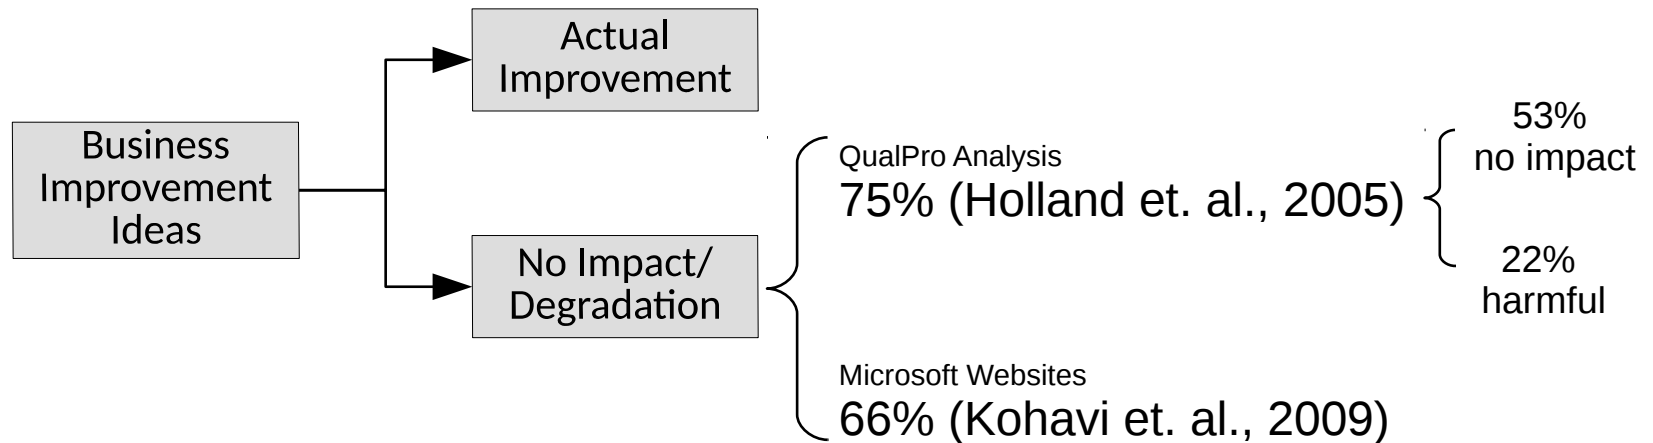


Outline



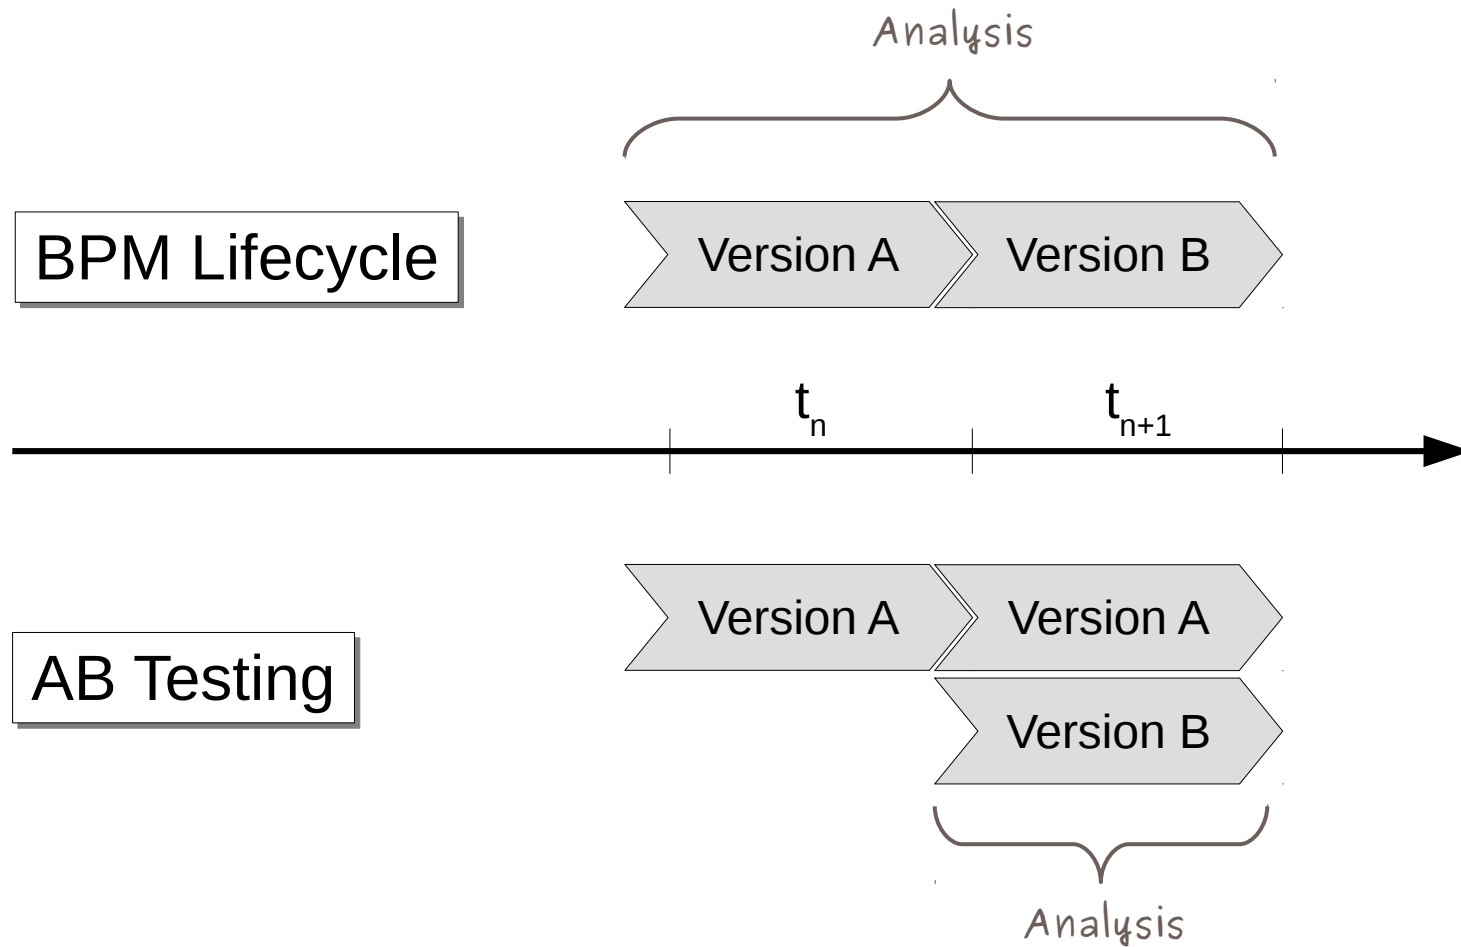
- Motivation
- AB Testing, AB-BPM approach, and their limitations
- Requirements
- Proposed Solution
 - System Architecture
 - ProcessBandit Algorithm
 - Reward Design
- Evaluation
- Conclusion

Motivation



- Business Process
 - Chain of Events, Activities, and Decisions
 - Expressed as Process Models
 - Instantiated and Executed by a Process Engine

Improvement Approach




AB Testing

- Randomized Experiment in production – A vs B vs ...
- Used to test micro-changes in web applications
- Test fairly, fail fast
- What if the test is risky?

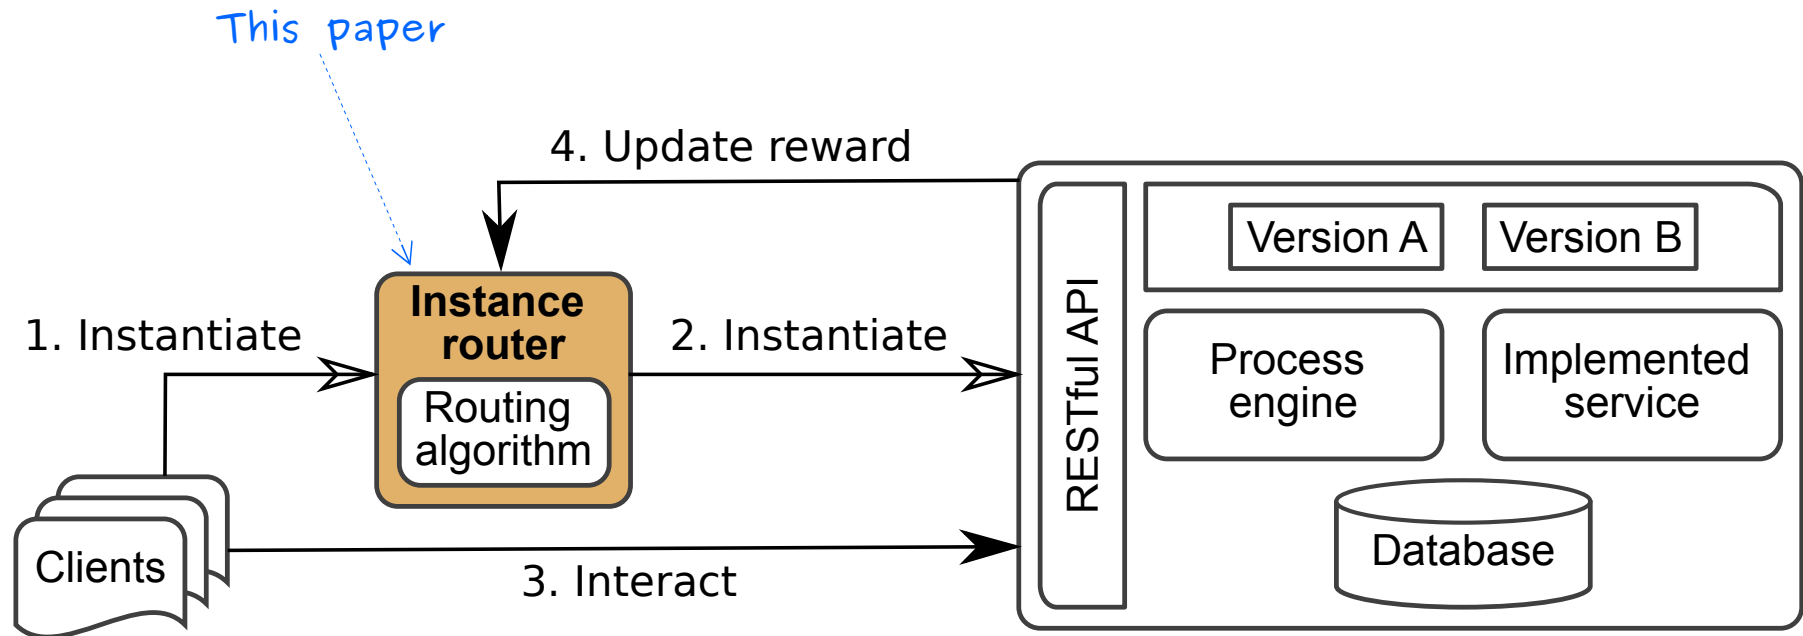
Example of risk management
Thompson Sampling in Clinical trials

↓
Routing algorithm



Icons made by Freepik from www.flaticon.com


AB-BPM Architecture



Source: Satyal et. al. (2017)

Routing with multi-armed bandits


- Contextual multi-armed bandit
- Based on LinUCB (Li et. al. 2010)
- Arm = Process Version
- Phase 1 (P1) – Experimentation (or exploration)
 - Instantiate some process versions at random (decay fxn + LinUCB approach)
 - Observe rewards only for processes instantiated during this phase
- Phase 2 (P2) - Exploitation
 - Do not observe rewards



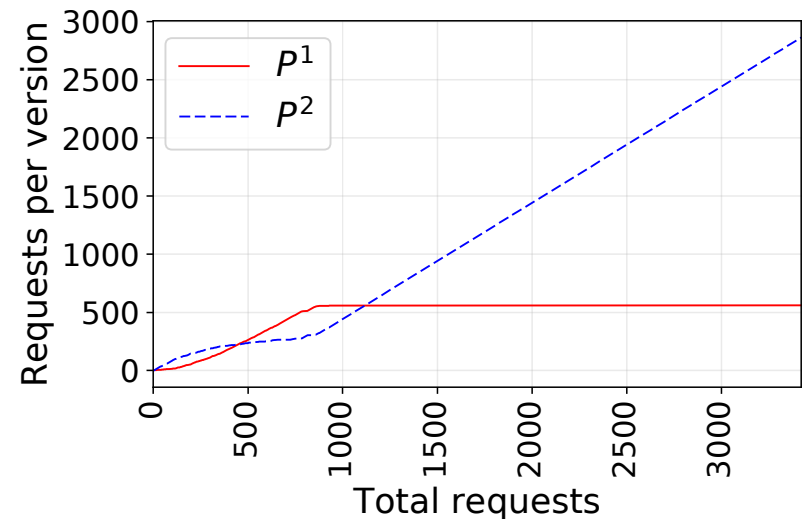
LTAvgR Algorithm and Rewards

- Designed for long running business processes
- One performance indicator (e.g. duration)

Good Rewards come early,
bad ones come late



Send more requests to
Better version over time




Source: Satyal et. al. (2017)


LTAvgR Algorithm and Rewards

- Designed for long running business processes
- One performance indicator (e.g. duration)

Good Rewards come early,
bad ones come late



Send more requests to
Better version over time



What about complex scenarios?

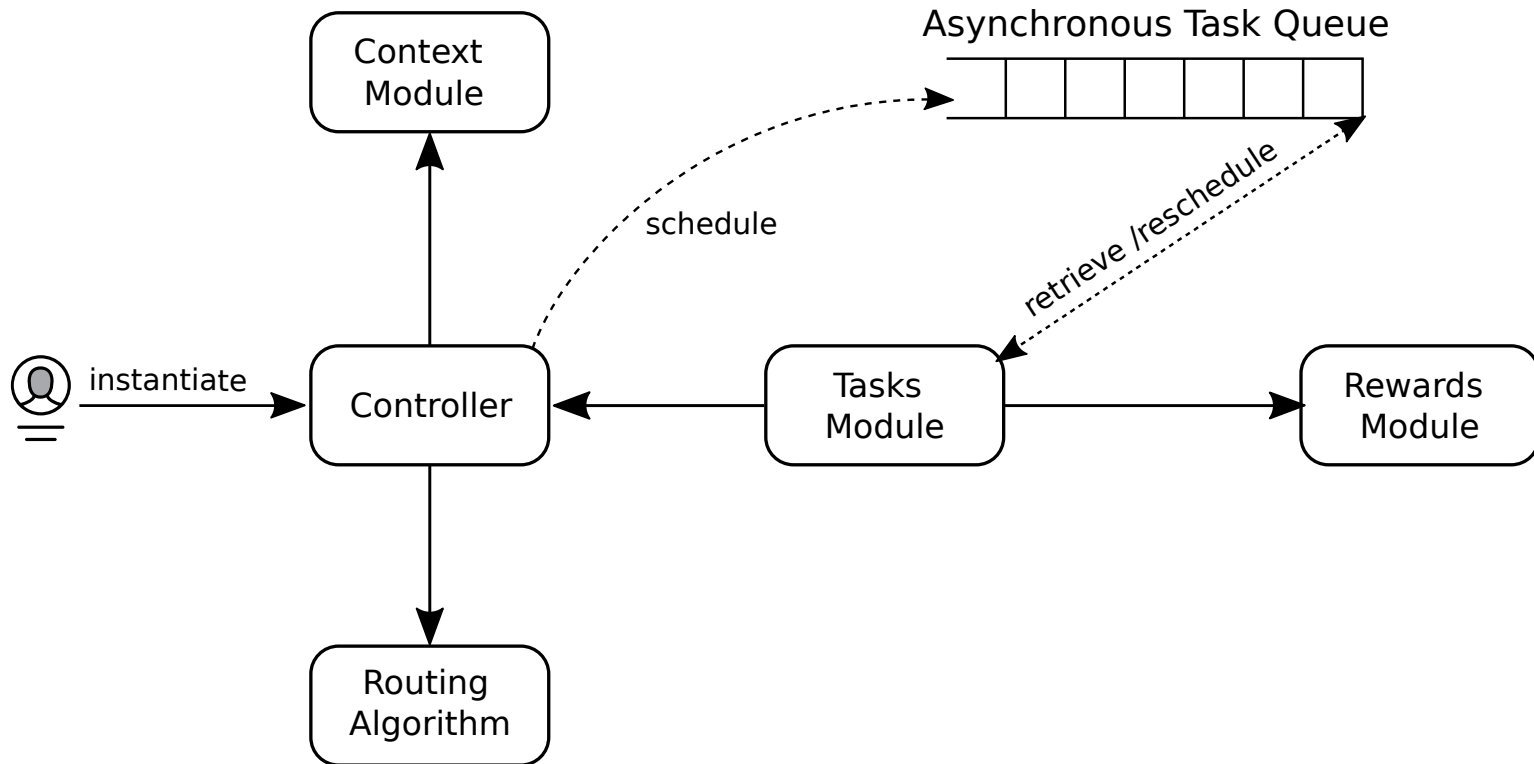
Source: Satyal et. al. (2017)

General BPM Scenarios



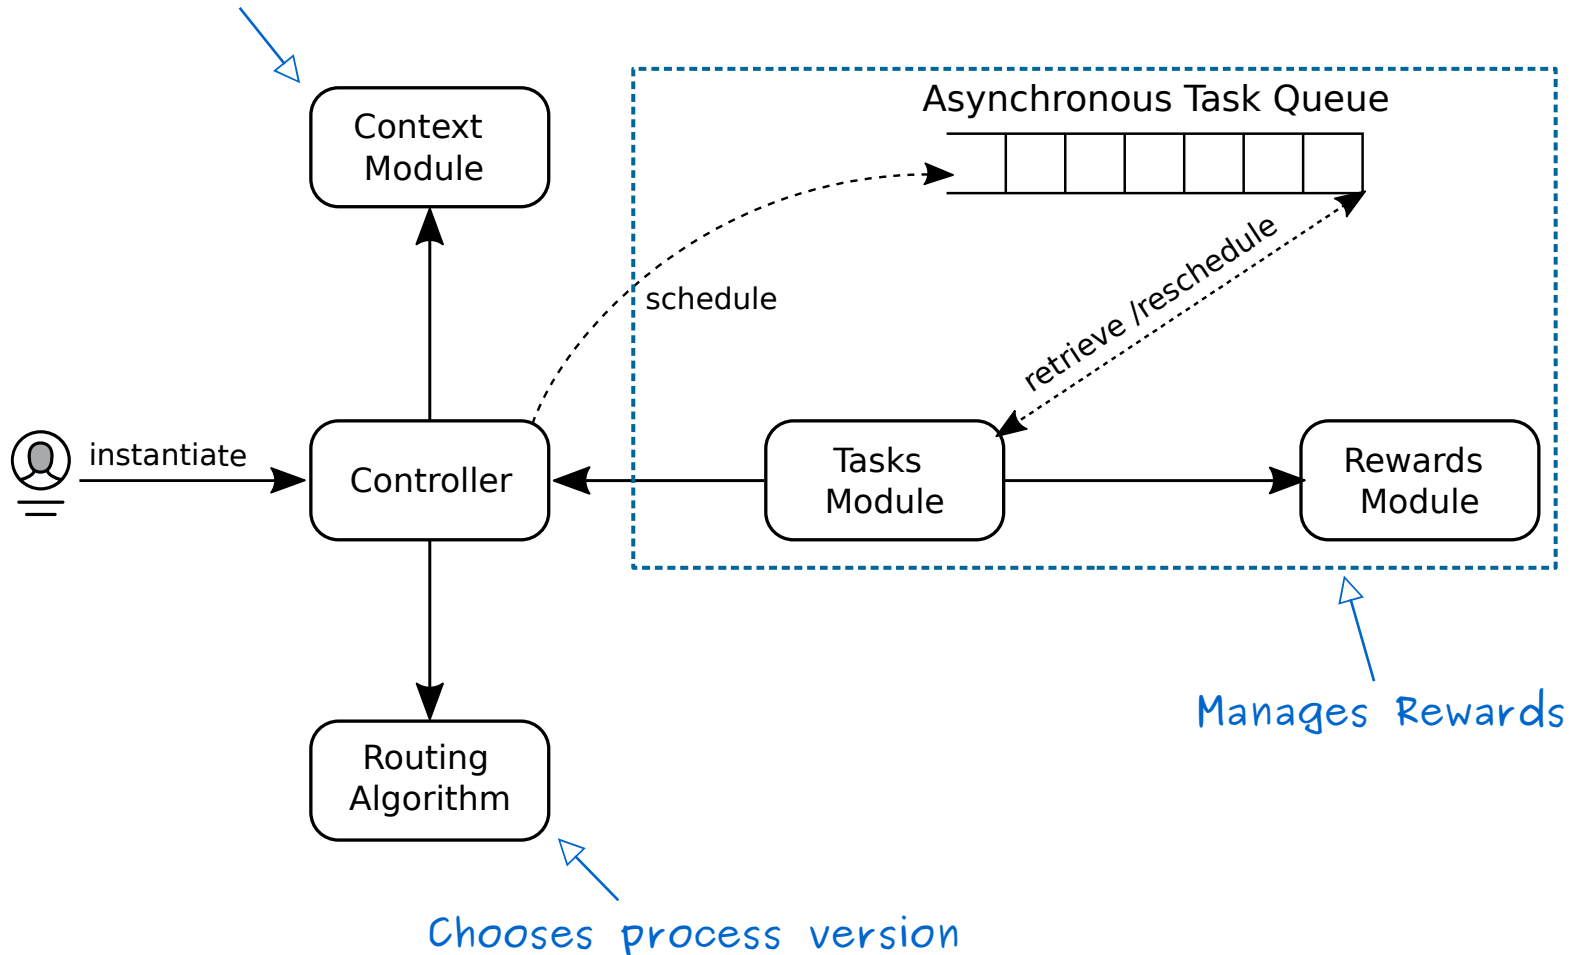
- Performance is determined by multiple Process Performance Indicators (PPIs) e.g. profit margin & user satisfaction
- Individual PPIs are available at different times
- Most process instances do not provide all PPIs e.g. user satisfaction
- Performance is affected by factors external to the process e.g. Weather Conditions (van der Aalst, et. al. 2007)

Instance Router Architecture




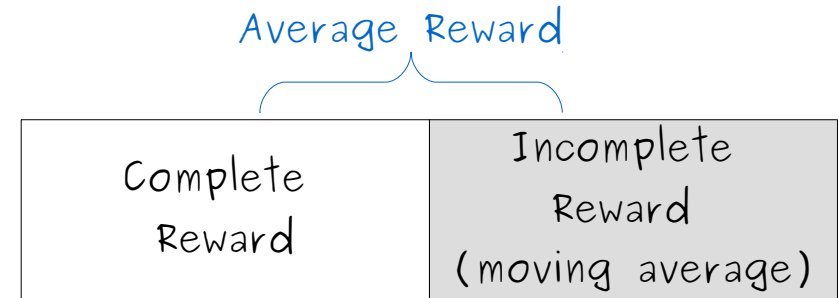
Instance Router Architecture

Identifies ext. contextual factors

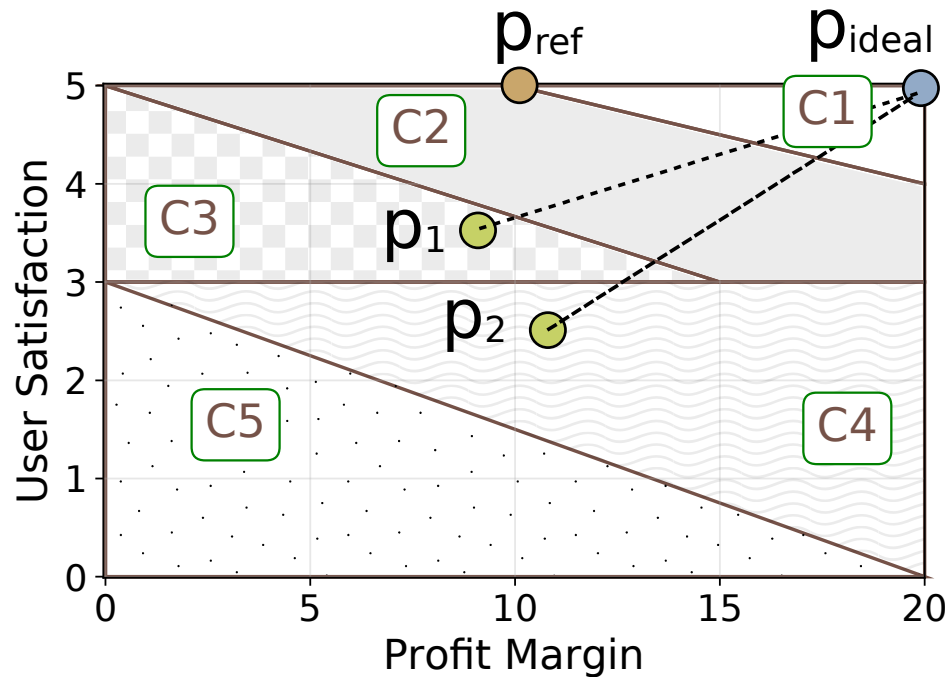


ProcessBandit Algorithm

- Fetch reward asynchronously
- Apply partial rewards
- Limit no. of incomplete rewards
- 3 Phased approach
 - P1A** Collect external contextual data, use only request information
 - P1B** Reset algorithm & use contextual data (if necessary)
 - P2** Send requests to version with best performance



Reward Design



$P_i = (\text{pr. margin, user sat.})$

$P_{\text{ideal}} = (\text{best pr. margin, best user sat.})$

$P_{\text{ref}} = (\text{assumed pr. margin, assumed user sat.})$

Distance based – implicit, too fine grained
What does dist. P1 vs dist. P2 indicate?

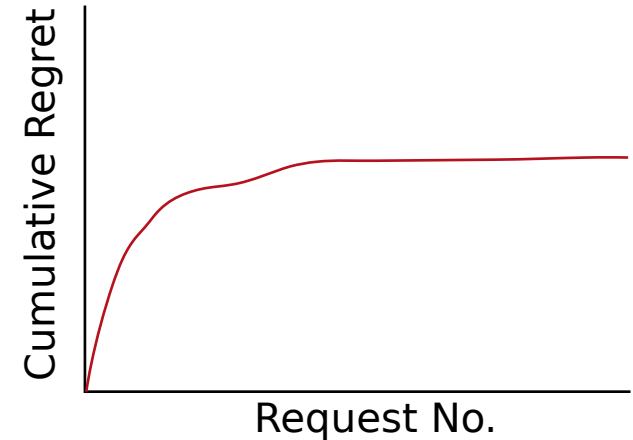
✗

Classification based – Explicit, coarse grained
e.g. C1 is x times better than C2

✓

Experiment Setup

- Regret
 - reward from best version – observed reward
 - Zero-regret strategy
 - avg. regret per request tends to 0



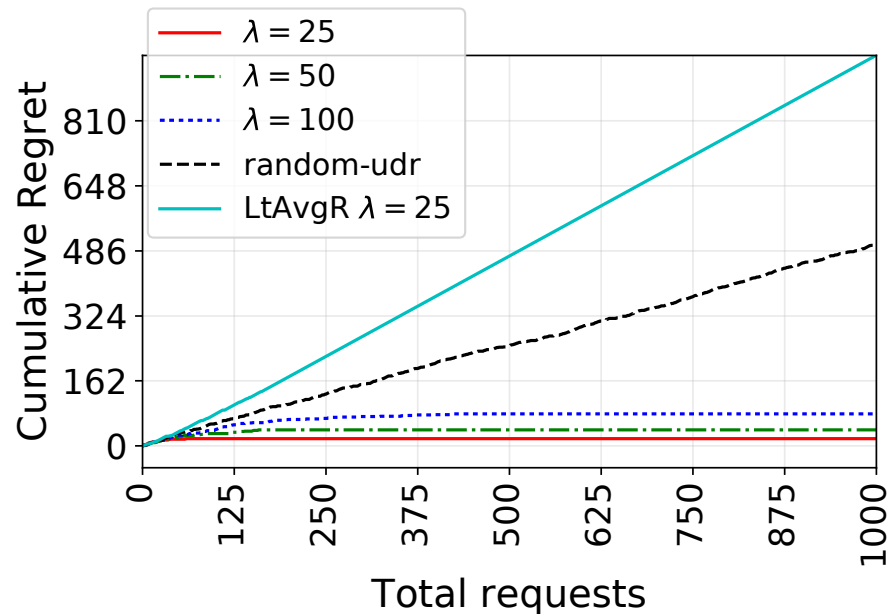
- Best version can be found only if all PPIs, request context, and external contextual factors are available

Req.	Context	Contextual Factor $f=1$				Contextual Factor $f=2$			
		Profit Margin		User Satisfaction		Profit Margin		User Satisfaction	
		Version A	Version B	Version A	Version B	Version A	Version B	Version A	Version B
X		9	11	3	2.5	11	9	2.5	3
Y		11	9	2.5	3	9	11	3	2.5


- Distributed Application: Python + Nginx + Redis + Docker

Convergence Characteristics

All PPIs are available

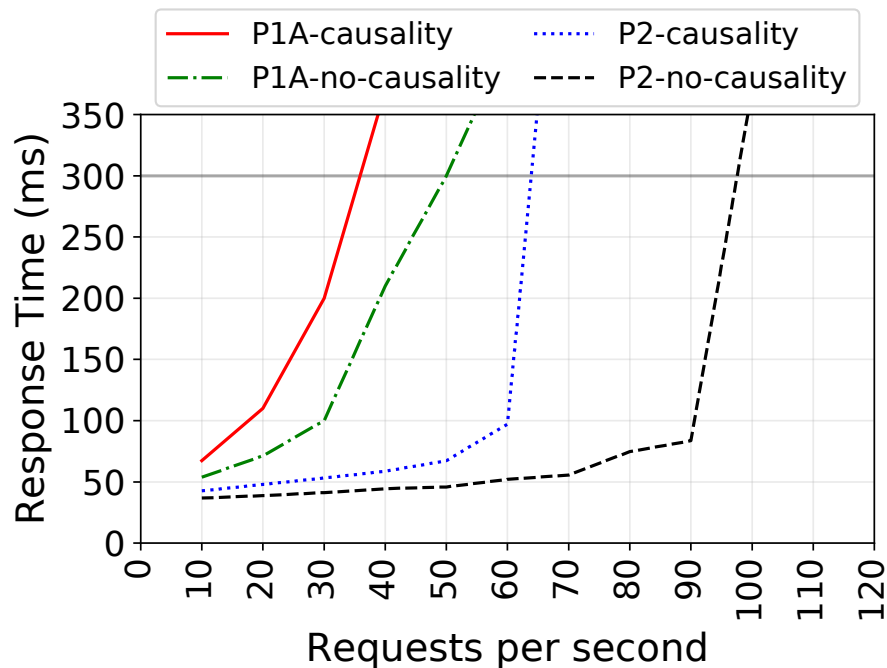


All PPIs are available
Only for $p\%$ instances




Scalability

Throughput Phase-Configuration

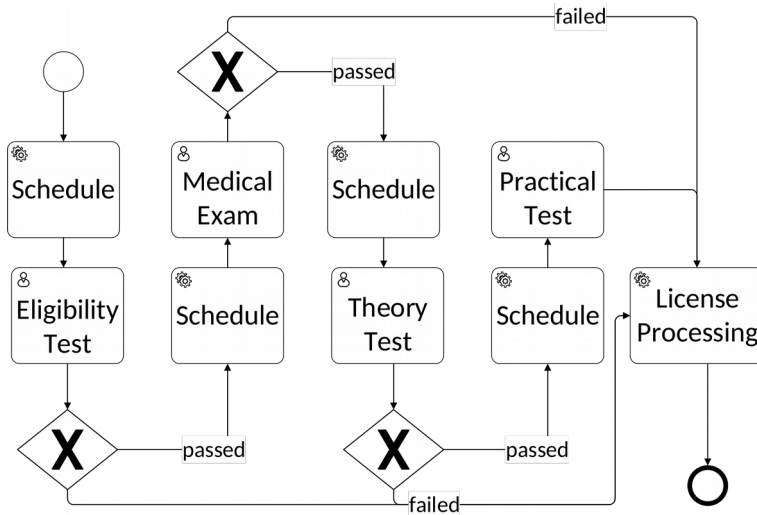


Horizontal Scalability i = alg. Server, r = redis server

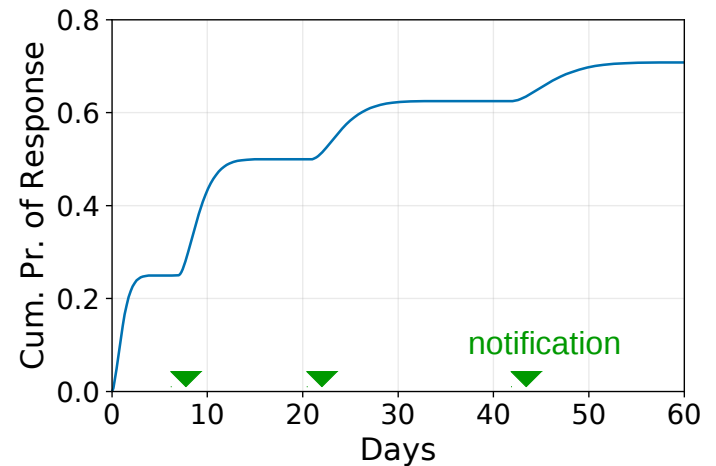

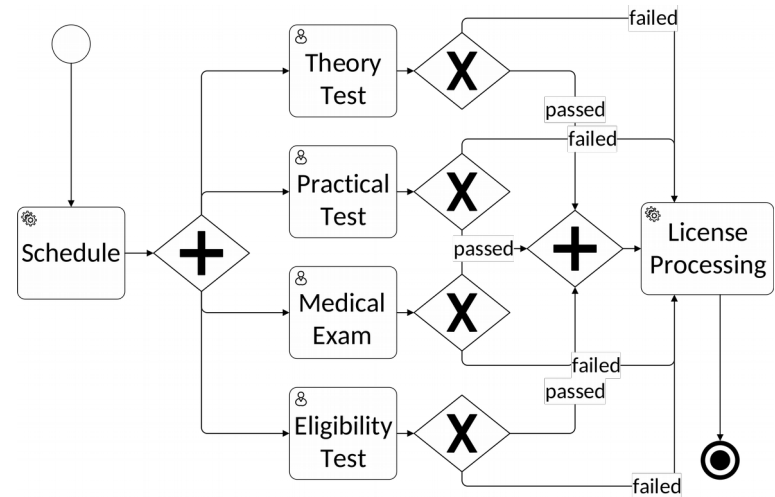


Pilot Licensing Scenario

Version A



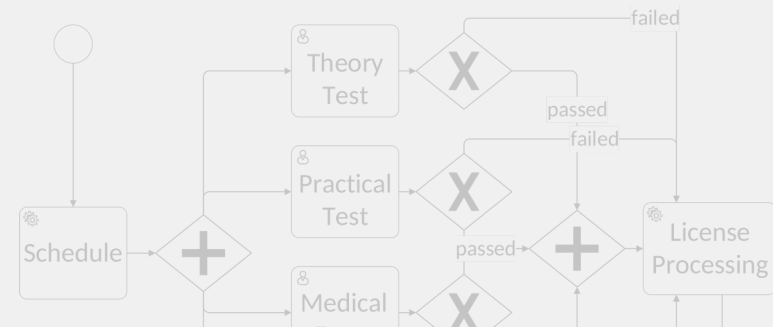
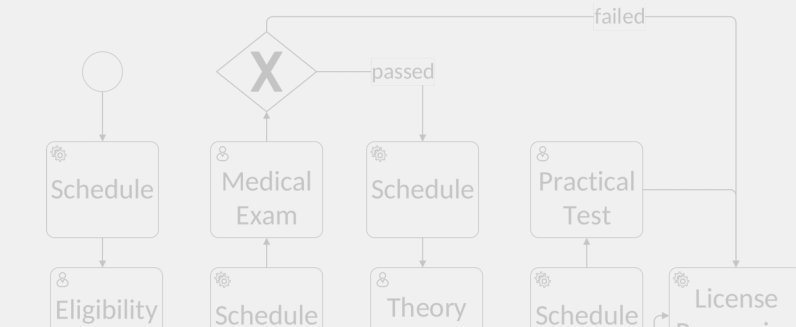
Version B





Pilot Licensing Scenario

Version A

Version B



ProcessBandit sends more requests to the version that is better on average in the given context.



Conclusion



- AB Testing system and algorithm for generalized BPM scenarios
 - Modular architecture
 - ProcessBandit algorithm
- Part of AB-BPM approach for business process improvement
- Extension of our previous work (BPM 2017)
- Future work – Real world studies, Extension of AB-BPM methodology



DATA
61



Researchgate Project

Thank you

Suhrid Satyal

PhD Candidate

Architecture and Analytics Platforms Team (AAP), Data61
Computer Science and Engineering, UNSW

e Suhrid.Satyal@data61.csiro.au

Australian Technology Park, Lvl 5, 13 Garden Street,
Eveleigh, NSW 2015 | Locked Bag 9013, Alexandria NSW 2015
www.data61.csiro.au

www.csiro.au

